

Yuhan Chi

Undergraduate, School of Mathematical Sciences, Fudan University
yhchi25@m.fudan.edu.cn | chi-shan0707.github.io | github.com/Chi-Shan0707 | Online CV

PROFILE

Undergraduate at **Fudan University** (School of Mathematical Sciences), pursuing a double degree in **Information and Computing Science** and **Artificial Intelligence**.
GPA: **3.96/4.00** (Fall 2025), **4.00/4.00** (Spring 2026); Cumulative: **3.98/4.00**; Rank: **3/213** in cohort.

EDUCATION

Fudan University Shanghai, China
B.S. Candidate in Information and Computing Science & Artificial Intelligence (Double Degree), School of Mathematical Sciences.Sep. 2025 – Present

RESEARCH INTERESTS

- **Mathematical Interpretability and AI Trust:** grounding claims about model reliability in inspectable mathematical and statistical evidence
- **Reasoning Verification and Evaluation:** studying when trace-level signals are diagnostic, misleading, protocol-dependent, or domain-specific
- **Optimization, Control, and Sequential Decision-Making:** using mathematical tools to reason about uncertainty, feedback, constraints, and interaction with environments

SELECTED OPEN-SOURCE PROJECTS

token-verification-mirage [Paper](#) · [AI4Math Workshop](#)
• Poster, **ICML 2026 Workshop on AI for Math (AI4Math)**; solo-authored, full pipeline (dataset selection, generation, evaluation design, analysis, writing)
• Evaluation-protocol choices alone (pooling, in-sample scoring, direction-agnostic AUROC) shift apparent verification AUROC by up to **0.18**
• Under corrected within-problem, leave-one-run-out, fixed-direction scoring, shallow token statistics cluster at only **0.60–0.75 AU-ROC**; final-token entropy drops from 0.72–0.75 to 0.47–0.48 once the direction-agnostic reporting artifact is removed

code-not-text [Demo](#)
• Cross-domain measurement study across 31,040 math, science, and coding runs: AoA 0.958 (math), 0.799 (science), 0.434 (coding, below the random-direction baseline)
• Rules out five alternative explanations for the coding failure (model capacity, label scarcity, surface noise, feature coverage, judge framing) — all converge to the same ceiling; frames the result as measurement non-invariance

TinyLoRA-GRPO-Coder [DeepWiki](#)
• Adapts *Learning to Reason in 13 Parameters* (Morris et al., 2026) to competitive programming: trains 32 shared scalars via GRPO on Qwen2.5-Coder-3B, rewarded by real g++ compile-and-run outcomes

microgpt.cpp
• Minimal GPT (autograd, multi-head attention, Adam) from scratch in ~300 lines of C++, inspired by Karpathy’s teaching gist

ACADEMIC SERVICE

Reviewer, [ICML 2026 Workshop on AI for Math \(AI4Math\)](#), 2026
Talk: **Reinforcement Learning: From Bandits to PPO** — *Apr 18, 2026* ([PDF notes](#))
• Overview covering multi-armed bandits, MDPs, policy gradient, and PPO

SKILLS

Programming	Python, C++, C
ML/AI	PyTorch, ML experimentation, LLM evaluation, RL basics
Tools	Git, GitHub, Linux, LaTeX, Markdown
Language	Chinese(native) English(fluent)

SELECTED COURSE GRADES

Semester	Course	Grade	Semester	Course	Grade
Fall 2025	Programming	A	Spring 2026	Mathematical Analysis II	A+
	Analytic Geometry	A		Advanced Algebra II	A
	Mathematical Analysis I	A		Foundations of Software for AI	A
	Advanced Algebra I	A-		Introduction to AI	A

COMMUNITY INVOLVEMENT

- [github-unflag-playbook-cn](#) — Chinese playbook for GitHub account flagging/recovery

- **ic-guide** — Open-source self-learning guide for integrated circuits
- **FDUGuideBook/nav-site** — Student navigation site for the Fudan community
- **FDU-Sharing** — Mutual-aid course-material sharing
- **Fudan Open Source Initiative (FDU-OSI)** — Founder; an initiative to promote the establishment and exchange of Fudan's open-source community through standardization